

# Bootcamp Project 2 – Financial Data Analysis

## Outline of the Project:

You are given 5 csv files which represent data related to Transactions and Loan acquired by Customers. Your task is to do proper ETL process for the data which includes data cleaning and transformation for storage purpose and store the processed data into SCD Type 1 Dimension Tables. Also, create table which encapsulate combined data created by joining raw tables. And At last, populate the data to Desktop Power BI.

Use only Databricks for all the computation and transformations.

## Work Flow of Project:





### Storing of the CSV files in ADLS G2 (bronze layer)

- I've created 3 directories in ADLS G2 (bronze, silver, Gold) which will be used in the project to follow medilean structure:

**Authentication method:** Access key ([Switch to Microsoft Entra user account](#))

**Location:** [raw](#) / Project-2

Search blobs by prefix (case-sensitive)







	Name	Modified	Access tier	Archive status	Blob type
<input type="checkbox"/>	 [..]				
<input type="checkbox"/>	 Bronze	4/24/2025, 9:09:54 PM			
<input type="checkbox"/>	 Gold	4/24/2025, 9:10:07 PM			
<input type="checkbox"/>	 Silver	4/24/2025, 9:10:00 PM			

- I've stored 5 csv files in bronze directory:

**Authentication method:** Access key ([Switch to Microsoft Entra user account](#))




**Location:** [raw](#) / [Project-2](#) / Bronze

Search blobs by prefix (case-sensitive)

Name	Modified	Access tier	Archive status	Blob type
<input type="checkbox"/>  [..]				
<input type="checkbox"/>  accounts.csv	4/27/2025, 8:36:54 AM	Hot (Inferred)		Block blob
<input type="checkbox"/>  customers.csv	4/27/2025, 8:36:54 AM	Hot (Inferred)		Block blob
<input type="checkbox"/>  loan_payments.csv	4/27/2025, 8:36:54 AM	Hot (Inferred)		Block blob
<input type="checkbox"/>  loans.csv	4/27/2025, 8:36:54 AM	Hot (Inferred)		Block blob
<input type="checkbox"/>  transactions.csv	4/27/2025, 8:36:54 AM	Hot (Inferred)		Block blob

## Creating connection between ADLS G2 and databricks using Service principal method:

- To create service principal, we need to do new app registration in Microsoft Entra ID:

 Delete  Endpoints  Preview features

### ^ Essentials

Display name	: <a href="#">new-azure-databricks</a>	Client credentials	: <a href="#">0 certificate, 1 secret</a>
Application (client) ID	: 2e763201-bf64-4120-834b-70c9b85cb1e7	Redirect URIs	: <a href="#">Add a Redirect URI</a>
Object ID	: a510edb2-939e-4b3a-9497-19b0c7260031	Application ID URI	: <a href="#">Add an Application ID URI</a>
Directory (tenant) ID	: 7a0f9522-2fed-4e78-aef0-64d397473d8d	Managed application in I...	: <a href="#">new-azure-databricks</a>
Supported account types	: <a href="#">My organization only</a>		

Store application ID and Tenant ID of the app in your system for future use.

Next, we need to create Client secret.

Create the secret and store it's value with you for future use (we can't access this value later)

After saving all 3 info in local, we need to give Storage Blob contributor role to this new service principal.

To do that, go to ADLS G2 container> IAM access> add new role>select Blob Storage contributor role> assign to new service principal name.

This screenshot shows the 'Access Control (IAM)' page for a resource. It displays role assignments for the 'Storage Blob Data Contributor' role. The table lists users and service principals with their respective permissions and subscriptions.

Role	Principal	Permissions	Subscription
Owner	Jassi Singh (User)	Owner	Subscription (Inherited)
	Ustaad (User)	Owner	Subscription (Inherited)
Storage Blob Data Contributor	azure-databricks (Service principal)	Storage Blob Data Contributor	This resource
	Jassi Singh (User)	Storage Blob Data Contributor	This resource
	new-azure-databricks (Service principal)	Storage Blob Data Contributor	This resource
	synapsewsp-1 (Managed identity)	Storage Blob Data Contributor	This resource

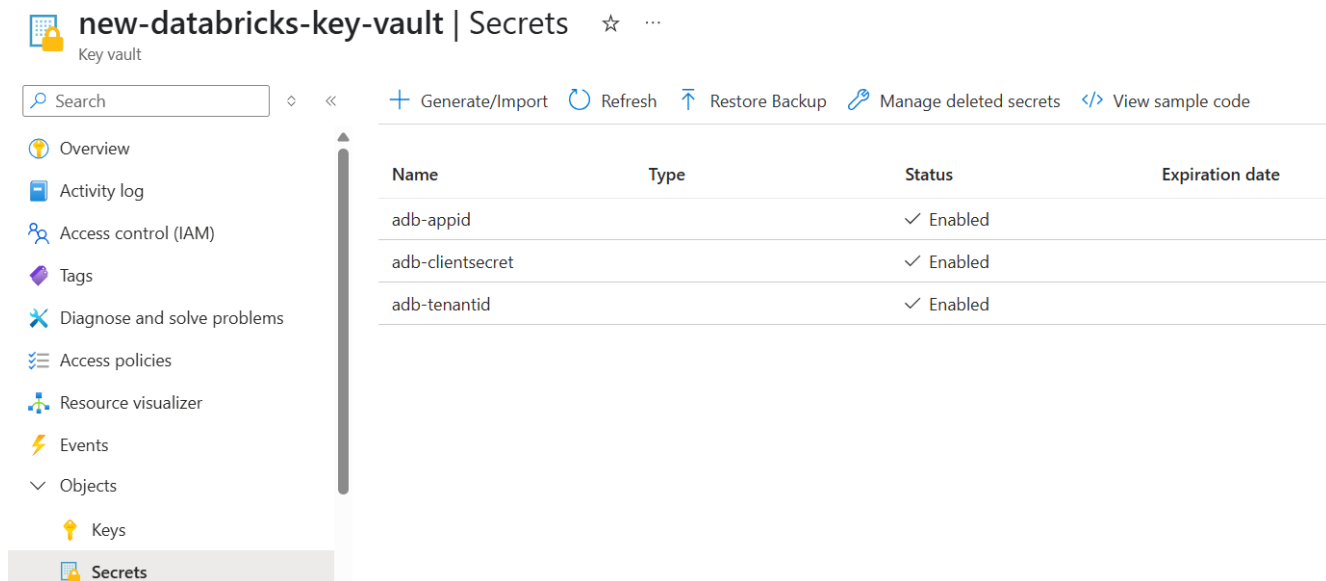
- Creation of key vault and add 3 secrets in key vault and give permissions of reading and listing secrets to New service principal:

This screenshot shows the 'Key vaults' page in the Azure portal. It displays a list of key vaults with their names, types, resource groups, locations, and subscriptions.

Name	Type	Resource Group	Location	Subscription
dev-key-vault-mine	Key vault	dev	East US	Azure subscription 1
new-databricks-key-vault	Key vault	dev	East US	Azure subscription 1

This screenshot shows the 'Access policies' page for the 'new-databricks-key-vault'. It displays a list of access policies with their names, emails, and permissions.

Name	Email	Key Permissions	Secret Permissions	Certificate Permissions
<b>APPLICATION</b>				
AzureDatabricks			Get, List	
new-azure-databricks			Get, List	
<b>USER</b>				
Jassi Singh	jassi_mann13_outlook.co...	Get, List, Update, Create, ...	Get, List, Set, Delete, Rec...	Get, List, Update,



new-databricks-key-vault | Secrets ☆ ...

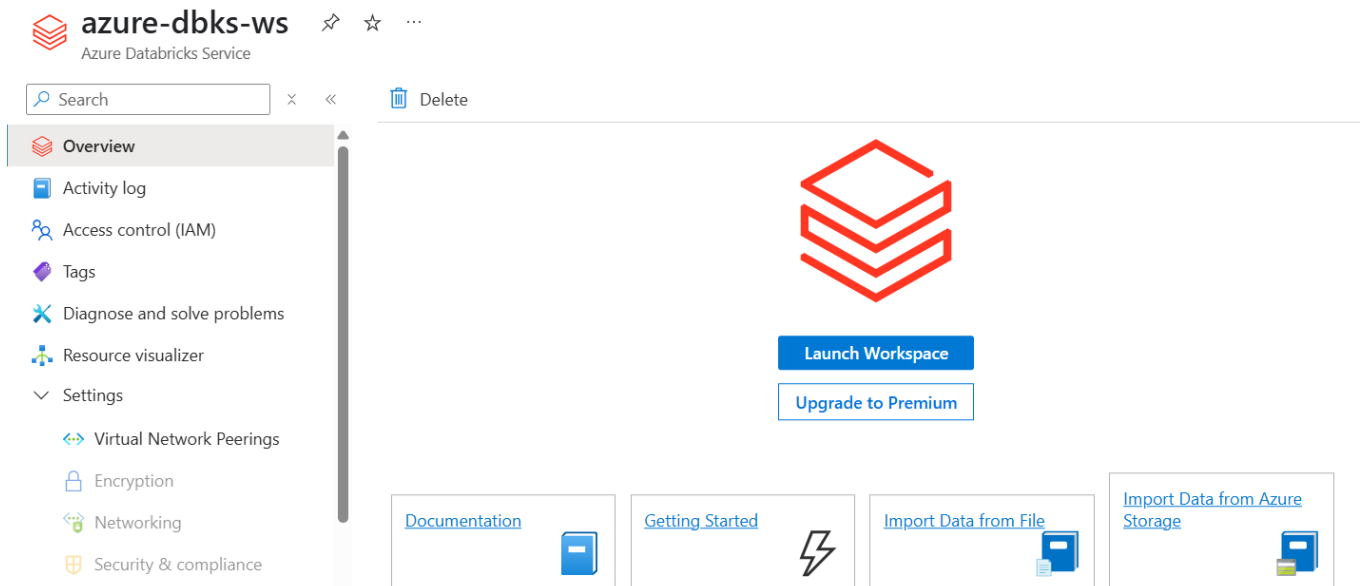
Key vault

Search

+ Generate/Import Refresh Restore Backup Manage deleted secrets View sample code

Name	Type	Status	Expiration date
adb-appid		✓ Enabled	
adb-clientsecret		✓ Enabled	
adb-tenantid		✓ Enabled	

- Creation of Azure Databricks Workspace



azure-dbks-ws ☆ ☆ ...

Azure Databricks Service

Search Delete

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Resource visualizer

Settings

Virtual Network Peerings

Encryption

Networking

Security & compliance

Launch Workspace

Upgrade to Premium

Documentation

Getting Started

Import Data from File

Import Data from Azure Storage

- To access ADLS G2 from Databricks notebook, we need to create scope in notebook so that, we can access key vault secrets we have created. For that, we need connection string/URI of Key vault and resource ID of it too.
- Go to <-----databricks url----->/#secrets/createScope
- Then you can access the key vault secrets by using dbutils.secrets commands

```
▶ ✓ Yesterday (1s) 4  
dbutils.secrets.list('scope-to-new-keyvault')  
[SecretMetadata(key='adb-appid'),  
 SecretMetadata(key='adb-clientsecret'),  
 SecretMetadata(key='adb-tenantid')]
```

- Connect with ADLS G2 using service principal method by using following Code:

```

storage_account_name = "azuredlstorageaccount" # e.g., "mystorageaccount"
container_name = "raw" # e.g., "mycontainer"
scope_name = "scope-to-new-keyvault" # e.g., "kvi-test"

# Retrieve secrets from Key Vault
application_id = dbutils.secrets.get(scope=scope_name, key="adb-appid")
client_secret = dbutils.secrets.get(scope=scope_name, key="adb-clientsecret")
tenant_id = dbutils.secrets.get(scope=scope_name, key="adb-tenantid")
# Define mount point
mount_point = "/mnt/project2" # e.g., "/mnt/mydata"

configs = {
    f"fs.azure.account.auth.type": "OAuth",
    f"fs.azure.account.oauth.provider.type": "org.apache.hadoop.fs.azurebfs.oauth2.ClientCredsTokenProvider",
    f"fs.azure.account.oauth2.client.id": application_id,
    f"fs.azure.account.oauth2.client.secret": client_secret,
    f"fs.azure.account.oauth2.client.endpoint": f"https://login.microsoftonline.com/{tenant_id}/oauth2/token"
}

# Check if a specific mount point exists
# mount_point = "/mnt/my-data"
mounts = dbutils.fs.mounts()

is_mounted = any(mount.mountPoint == mount_point for mount in mounts)

if is_mounted:
    print(f"Mount point {mount_point} already exists.")
else:
    print(f"Mount point {mount_point} does not exist and creating new mount point")
    dbutils.fs.mount(
        source="abfss://raw@azuredlstorageaccount.dfs.core.windows.net/Project-2",
        mount_point="/mnt/project2",
        extra_configs=configs
    )

# List files in the mount point
dbutils.fs.ls(mount_point)

```

Mount point /mnt/project2 already exists.

```

[FileInfo(path='dbfs:/mnt/project2/Bronze/', name='Bronze/', size=0, modificationTime=1745543394000),
 FileInfo(path='dbfs:/mnt/project2/Gold/', name='Gold/', size=0, modificationTime=1745543407000),
 FileInfo(path='dbfs:/mnt/project2/Silver/', name='Silver/', size=0, modificationTime=1745543400000)]

```

- To Read CSV file from ADLS G2 directory :

```
▶ Yesterday (17s) 8

df_accounts = spark.read.csv("/mnt/project2/Bronze/accounts.csv", header=True,inferSchema=True)
df_accounts.printSchema()
df_accounts.show()

▶ (3) Spark Jobs

▶ df_accounts: pyspark.sql.dataframe.DataFrame = [account_id: integer, customer_id: integer ... 2 more fields]
|      3|      78|    Savings| 1500.0|
|      4|      34|   Checking|3000.25|
|      5|      56|    Savings|   500.0|
|      6|      23|   Checking| 1200.5|
|      7|      89|    Savings|  800.75|
|      8|      67|   Checking| 2200.0|
|      9|      14|    Savings|   900.25|
|     10|      92|   Checking| 1800.5|
|     11|       3|    Savings|1100.75|
|     12|      81|   Checking| 2700.0|
|     13|      29|    Savings|1300.25|
```

- Data Transformation I did were :

```
▶ Yesterday (4s) 9

from pyspark.sql.functions import *
df_accounts = df_accounts.dropDuplicates()
# -----
df_accounts = df_accounts.where(col("account_id").isNotNull() & col("customer_id").isNotNull())
# -----
df_accounts= df_accounts.fillna(0.0,"balance")
df_accounts.show()

▶ (2) Spark Jobs

▶ df_accounts: pyspark.sql.dataframe.DataFrame = [account_id: integer, customer_id: integer ... 2 more fields]
|      4|      34|   Checking|3000.25|
|      2|      12|   Checking|2500.75|
|      1|      45|    Savings| 1000.5|
|      8|      67|   Checking| 2200.0|
|     21|      53|    Savings|   300.25|
|      6|      23|   Checking| 1200.5|
|     17|      90|    Savings|   600.25|
```

1. Dropping duplicates records
2. Filtering out records, which have non-null values of particular columns
3. Setting up default values to those column values which will be NULL.

- I've applied these transformations to all 5 Dataframes:

```

▶ ✓ Yesterday (1s) 12 Python
from pyspark.sql.functions import *
df_customers = df_customers.dropDuplicates()
# -----
df_customers = df_customers.where(col("customer_id").isNotNull() & col("first_name").isNotNull() & col("last_name").isNotNull())
df_customers.show()

▶ (2) Spark Jobs

▶ df_customers: pyspark.sql.dataframe.DataFrame = [customer_id: integer, first_name: string ... 5 more fields]
| 11| Alexander| Thomas| 1010 Willow Rd| St. John's| NL| A1A0A1|
| 4| Emily| Davis| 101 Pine Rd| Calgary| AB| T2A0A1|
| 21| Andrew| Mitchell| 2020 Spruce Ln| Hamilton| ON| L8P0A1|
| 1| John| Doe| 123 Elm St| Toronto| ON| M4B1B3|
| 6| Emma| Clark| 505 Cedar St| Halifax| NS| B2H0A1|

```

```

▶ ✓ Yesterday (1s) 15
from pyspark.sql.functions import *
df_loanpayments = df_loanpayments.dropDuplicates()
# -----
df_loanpayments = df_loanpayments.where(col("payment_id").isNotNull() & col("loan_id").isNotNull())

df_loanpayments = df_loanpayments.fillna(0.0, "payment_amount")
df_loanpayments.show()

▶ (2) Spark Jobs

▶ df_loanpayments: pyspark.sql.dataframe.DataFrame = [payment_id: integer, loan_id: integer ... 2 more fields]
| 19| 10| 2024-01-19| 1000.0|
| 8| 78| 2024-01-08| 450.0|
| 18| 99| 2024-01-18| 950.0|
| 14| 55| 2024-01-14| 750.0|
| 7| 56| 2024-01-07| 400.0|
| 6| 34| 2024-01-06| 350.0|

```

```

▶ ✓ Yesterday (2s) 18
from pyspark.sql.functions import *
df_loans = df_loans.dropDuplicates()
# -----
df_loans = df_loans.where(col("loan_id").isNotNull() & col("customer_id").isNotNull())

df_loans = df_loans.fillna(0.0, ["loan_amount", "interest_rate"])
df_loans = df_loans.fillna(0, "loan_term")
df_loans.show()

▶ (2) Spark Jobs

▶ df_loans: pyspark.sql.dataframe.DataFrame = [loan_id: integer, customer_id: integer ... 3 more fields]
| 17| 99| 22500.25| 5.5| 36|
| 10| 92| 37500.5| 4.5| 48|
| 12| 81| 20000.0| 3.5| 24|
| 3| 78| 15000.0| 6.0| 60|
| 20| 21| 37500.0| 3.5| 24|

```



```
▶ ✓ Yesterday (1s) 21

from pyspark.sql.functions import *
df_transactions = df_transactions.dropDuplicates()
# -----
df_transactions = df_transactions.where(col("transaction_id").isNotNull() & col("account_id").isNotNull())

df_transactions = df_transactions.fillna(0.00,"transaction_amount")
df_transactions.show()
```

▶ (2) Spark Jobs

▶ df\_transactions: pyspark.sql.dataframe.DataFrame = [transaction\_id: integer, account\_id: integer ... 3 more fields]

1	45	2024-01-01	100.5	Deposit
2	12	2024-01-02	200.75	Withdrawal
18	5	2024-01-18	275.75	Withdrawal
15	47	2024-01-15	250.0	Deposit
8	67	2024-01-08	275.75	Withdrawal
16	18	2024-01-16	175.0	Withdrawal

- Saving all transformed Data frames into delta format in silver layer – ADLS G2

```
▶ ✓ Yesterday (50s) 23

delta_path = "/mnt/project2/Silver/"

# Save DataFrame as Delta table
df_accounts.write \
    .format("delta") \
    .mode("overwrite") \
    .save(delta_path+"Accounts/")

# Save DataFrame as Delta table
df_customers.write \
    .format("delta") \
    .mode("overwrite") \
    .save(delta_path+"Customers/")

# Save DataFrame as Delta table
df_loans.write \
    .format("delta") \
    .mode("overwrite") \
    .save(delta_path+"Loans/")

# Save DataFrame as Delta table
● df_loanpayments.write \
    .format("delta") \
    .mode("overwrite") \
    .save(delta_path+"Loan-payments/")

# Save DataFrame as Delta table
df_transactions.write \
    .format("delta") \
    .mode("overwrite") \
    .save(delta_path+"Transactions/")

▶ (35) Spark Jobs
```

- To combine all Data frames into one common DF by using their relationships with each other, I've created a new Data frame and store it into delta table in ADLS Silver layer.

```

▶ ✓ Yesterday (13s) 24

df_acc = df_accounts.join(df_customers, on="customer_id", how="inner")

df_loans1= df_loans.join(df_loanpayments, on="loan_id", how="inner")

df_acc_new = df_acc.join(df_transactions, on="account_id", how="left")

df_final = df_acc_new.join(df_loans1, on="customer_id", how="left")

df_final = df_final.dropDuplicates()
df_final.show()

delta_path = "/mnt/project2/Silver/"
df_final.write \
    .format("delta") \
    .mode("overwrite") \
    .save(delta_path+"Curated-data/")

▶ (23) Spark Jobs

```

```

▶ (23) Spark Jobs

▶ df_acc: pyspark.sql.dataframe.DataFrame = [customer_id: integer, account_id: integer ... 8 more fields]
▶ df_acc_new: pyspark.sql.dataframe.DataFrame = [account_id: integer, customer_id: integer ... 12 more fields]
▶ df_final: pyspark.sql.dataframe.DataFrame = [customer_id: integer, account_id: integer ... 19 more fields]
▶ df_loans1: pyspark.sql.dataframe.DataFrame = [loan_id: integer, customer_id: integer ... 6 more fields]

```

024-01-16	175.0	Withdrawal	18	27500.5	4.5	48	47	2024-02-16	2400.0	▲
11	24	Checking	2600.0	Alexander	Thomas	1010 Willow Rd	St. John's	NL A1A0A1	46	2
024-02-15	175.0	Withdrawal	24	30000.0	3.0	24	93	2024-04-02	4700.0	
2	82	Checking	8300.5	Jane	Smith	456 Maple Ave	Ottawa	ON K1A0B1	83	2
024-03-23	150.0	Deposit	82	20000.5	4.5	48	71	2024-03-11	3600.0	
15	38	Checking	3900.5	Matthew	King	1414 Cedar Ln	Whitehorse	YT Y1A0A1	90	2
024-03-30	375.25	Withdrawal	38	27500.5	4.0	48	67	2024-03-07	3400.0	
13	44	Checking	4500.0	Daniel	Harris	1212 Ash Blvd	Charlottetown	PE C1A0A1	92	2
024-04-01	200.75	Withdrawal	44	30000.0	3.5	24	13	2024-01-13	700.0	
12	64	Checking	6500.0	Isabella	Lee	1111 Poplar St	Fredericton	NB E3B0A1	14	2
024-01-14	300.25	Withdrawal	2	20000.75	4.5	48	91	2024-03-31	4600.0	
10	52	Checking	5300.0	Ava	Anderson	909 Cypress Ave	Quebec City	QC G1A0A1	61	2
024-03-01	100.5	Deposit	52	20000.0	3.5	24	41	2024-02-10	2100.0	
14	9	Savings	900.25	Sophia	Young	1313 Beech Dr	Yellowknife	NT X1A0A1	32	2

## CONVERTING DATA INTO SCD TYPE 1 DIMENSION TABLES BY USING DATABRICKS:

```
▶ ✓ Yesterday (12s) 4 Python ✨ ⌵ ⋮

from pyspark.sql.functions import current_timestamp

# Define paths
source_path = "/mnt/project2/Silver/Accounts/"
target_path = "/mnt/project2/Gold/Accounts/"

# Step 1: Load source Delta table
try:
    source_df = spark.read.format("delta").load(source_path)
    source_df = source_df.withColumn("last_updated", current_timestamp())
    print("Source Delta table:")
    source_df.show(5)
    print(f"Source row count: {source_df.count()}")
except Exception as e:
    print(f"Error reading source Delta table: {str(e)}")
    raise

# Step 2: Load existing target Parquet (if exists)
try:
    target_df = spark.read.parquet(target_path)
    print("Existing Parquet file:")
    target_df.show(5)
    print(f"Target row count: {target_df.count()}")
except Exception as e:
    print(f"No existing Parquet file: {str(e)}")
    target_df = spark.createDataFrame([], source_df.schema)

# Step 3: Implement SCD Type 1
non_matching_target = target_df.join(source_df.select("account_id"), "account_id", "left_anti")
result_df = non_matching_target.unionByName(source_df)
print("Result DataFrame after SCD Type 1:")
result_df.show(5)
print(f"Result row count: {result_df.count()}")

# Step 4: Save as Parquet
try:
```

```

result_df.write \
    .format("parquet") \
    .mode("overwrite") \
    .option("compression", "snappy") \
    .save(target_path)
print(f"Parquet file saved to {target_path}")
except Exception as e:
    print(f"Error saving Parquet file: {str(e)}")

# Step 5: Verify
try:
    files = dbutils.fs.ls(target_path)
    print(f"Files in {target_path}: {files}")
    parquet_df = spark.read.parquet(target_path)
    print("Parquet file contents:")
    parquet_df.show(5)
    print(f"Parquet row count: {parquet_df.count()}")
except Exception as e:
    print(f"Error verifying Parquet file: {str(e)}")

```

► (16) Spark Jobs

- New Data frame created :

```

+-----+-----+-----+-----+-----+
|account_id|customer_id|account_type|balance|last_updated|
+-----+-----+-----+-----+-----+
|      36|      27|   Checking| 3700.0|2025-04-27 23:21:...|
|      85|      65|    Savings|  800.25|2025-04-27 23:21:...|
|      82|       2|   Checking| 8300.5|2025-04-27 23:21:...|
|      52|      10|   Checking| 5300.0|2025-04-27 23:21:...|
|      65|      69|    Savings|  550.25|2025-04-27 23:21:...|
+-----+-----+-----+-----+-----+
only showing top 5 rows

```

## Populating Data on Power BI Dashboard:

